

COMS 474
Homework 4

Haadi Majeed
Spring 2022

Problem 1

[24 points total (5,5,3,3,4,4)]

Suppose you use lasso fit to a linear model for a data set. Let $\beta^*(\lambda)$ denote the lasso solution for a specific λ (i.e. the coefficient vector you get for that λ).

Provide explanations for your answers to the following questions.

Part A


Describe how the training MSE changes as a function of λ , including $\lambda = 0$ and as $\lambda \rightarrow \infty$

The training MSE of the function $\beta^*(\lambda)$ would return the least squares fit when $\lambda = 0$. The value output does change though as λ grows in size, with the upper bound of ∞ being the most ideal constant that could be supplied to the equation, this is because at *some* point, λ will hit the absolute minimum of the function, which would be within the range of $0 \rightarrow \infty$.

Part B

Describe how the hold-out MSE changes as a function of λ , including $\lambda = 0$ and as $\lambda \rightarrow \infty$

The hold-out MSE *should* look similar to the training MSE and by applying a similar set of logic, the models should look relatively the same as it is from the same overall set of data. The resulting model should have an absolute minimum which would indicate what the most ideal λ would be. This value should be, if not exactly, then very close to where the training set's absolute minimum should be at. The image I have attached below is from the notes and displays what the models could look like as an example.



Part C

Describe $\beta^*(0)$.

When initially at $\beta^*(0)$, the function's value that it would return would simply be the L.S.F (least squares fit) value.

Part D

Describe what happens to $\beta^*(0)$ as λ grows.

The model $\beta^*(0)$ would cover every possible value between $0 \rightarrow \infty$ and would eventually hit the functions absolute minimum in a concave region. This is where it deviates off the null model and no longer just returns zero, and instead will indicate what the absolute minimum, or the most ideal λ , would be.

Part E

If you used ridge regression instead of lasso, explain how your answers to (a)-(d). would differ.

If we were to use ridge regression, we should expect a bit of variance in the data since they approach model slightly differently. Another attribute of ridge regression is that it is not capable to have sharp points in the model, thus intersections at an axis is not common. This subsequently results in non-zero estimations. The models generated by this when comparing the training vs the hold-out data should look very similar to each other due to data distribution. Ridge regression looks at both extremes and steps through the function looking at each step from $0 \rightarrow \infty$. It also penalises $\lambda \rightarrow \infty$ significantly more than it would $\lambda = 0$. When $\lambda = 0$ it would return the Least Squares estimation, while when $\lambda \rightarrow \infty$ it estimates the model will approach zero and the shrinkage penalty grows.

Part F

We discussed the “constrained form” of lasso, with a constraint of the form

$$\sum_{i=1}^p |\beta_j| \leq t$$

Which value, or limiting value, of t corresponds to $\lambda = 0$ and which corresponds to $\lambda \rightarrow \infty$

Problem 2

[15 points]

You have already seen formulas for the best intercept in linear models when there are no features $p = 0$ and a single feature $p = 1$. You will now look at what happens with p features when we center the data.

Recall that “centering” a feature means subtracting its mean. For example, if the sample

values for feature X_4 are $\begin{bmatrix} 5 \\ 0 \\ 1 \end{bmatrix}$, which has a mean of 2,

we could replace it with $\begin{bmatrix} 5-2 \\ 0-2 \\ 1-2 \end{bmatrix} = \begin{bmatrix} 3 \\ -2 \\ -1 \end{bmatrix}$ which has a mean of 0. Thus if feature X_4 is centred, then $\sum_{i=1}^n X_4(i) = 0$.

What is the value of the intercept β_0^* in the ordinary least squares solution, i.e.

$$(\beta_0^*, \beta_1^*, \dots, \beta_p^*) = \arg \min_{(\beta_0^*, \beta_1^*, \dots, \beta_p^*) \in \mathbb{R}^{p+1}} \frac{1}{n} \sum_{i=1}^n \left(Y(i) - \beta_0 - \sum_{j=1}^p \beta_j X_j(i) \right)^2$$

when the features X_1, \dots, X_p are all centered? (e.g. $\sum_{i=1}^n X_j(i) = 0$ for $j = 1, \dots, p$). (You do not need to use a second derivative test or solve for $\{\beta_1^*, \dots, \beta_p^*\}$, just use the first derivative test $0 = \frac{\partial}{\partial \beta_0} \text{MSE}$).